

Phase-1

✓ 1. Initial Data Review (Pandas)

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# URL provided by the user
url = "https://raw.githubusercontent.com/BidGitHub2022/AI-501-Final-Project/refs/heads/main/car_price_prediction_.csv"
df = pd.read_csv(url)

# 1. Basic Info
print("---- Dataset Info ----")
print(df.info())

# 2. Check for missing values
print("\n---- Missing Values ----")
print(df.isnull().sum())

# 3. Check for duplicates
print(f"\nDuplicate rows found: {df.duplicated().sum()}")
df = df.drop_duplicates()

# 4. High-level stats
print("\n---- Descriptive Statistics ----")
display(df.describe())

--- Dataset Info ---
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2500 entries, 0 to 2499
Data columns (total 10 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   Car ID      2500 non-null   int64  
 1   Brand        2500 non-null   object  
 2   Year         2500 non-null   int64  
 3   Engine Size 2500 non-null   float64 
 4   Fuel Type    2500 non-null   object  
 5   Transmission 2500 non-null   object  
 6   Mileage      2500 non-null   int64  
 7   Condition    2500 non-null   object  
 8   Price         2500 non-null   float64 
 9   Model         2500 non-null   object  
dtypes: float64(2), int64(3), object(5)
memory usage: 195.4+ KB
None

--- Missing Values ---
Car ID      0
Brand       0
Year        0
Engine Size 0
Fuel Type   0
Transmission 0
Mileage     0
Condition   0
Price       0
Model       0
dtype: int64

Duplicate rows found: 0

--- Descriptive Statistics ---
   Car ID      Year   Engine Size      Mileage      Price
count  2500.00000  2500.00000  2500.000000  2500.000000  2500.000000
mean   1250.50000  2011.6268   3.465240   149749.844800  52638.022532
std    721.83216   6.9917    1.432053   87919.952034  27295.833455
min    1.00000   2000.0000   1.000000   15.000000   5011.270000
25%    625.75000  2005.0000   2.200000   71831.500000  28908.485000
50%    1250.50000  2012.0000   3.400000   149085.000000  53485.240000
75%    1875.25000  2018.0000   4.700000   225990.500000  75838.532500
max    2500.00000  2023.0000   6.000000   299967.000000  99982.590000
```

Based on above results, dataset is remarkably clean—zero missing values and no duplicates.

Double-click (or enter) to edit

✓ 2. EDA

```
# Set visual style
sns.set_theme(style="whitegrid")

# 1. Price Distribution (Target Variable)
plt.figure(figsize=(10, 5))
sns.histplot(df['Price'], kde=True, bins=30, color='blue')
```

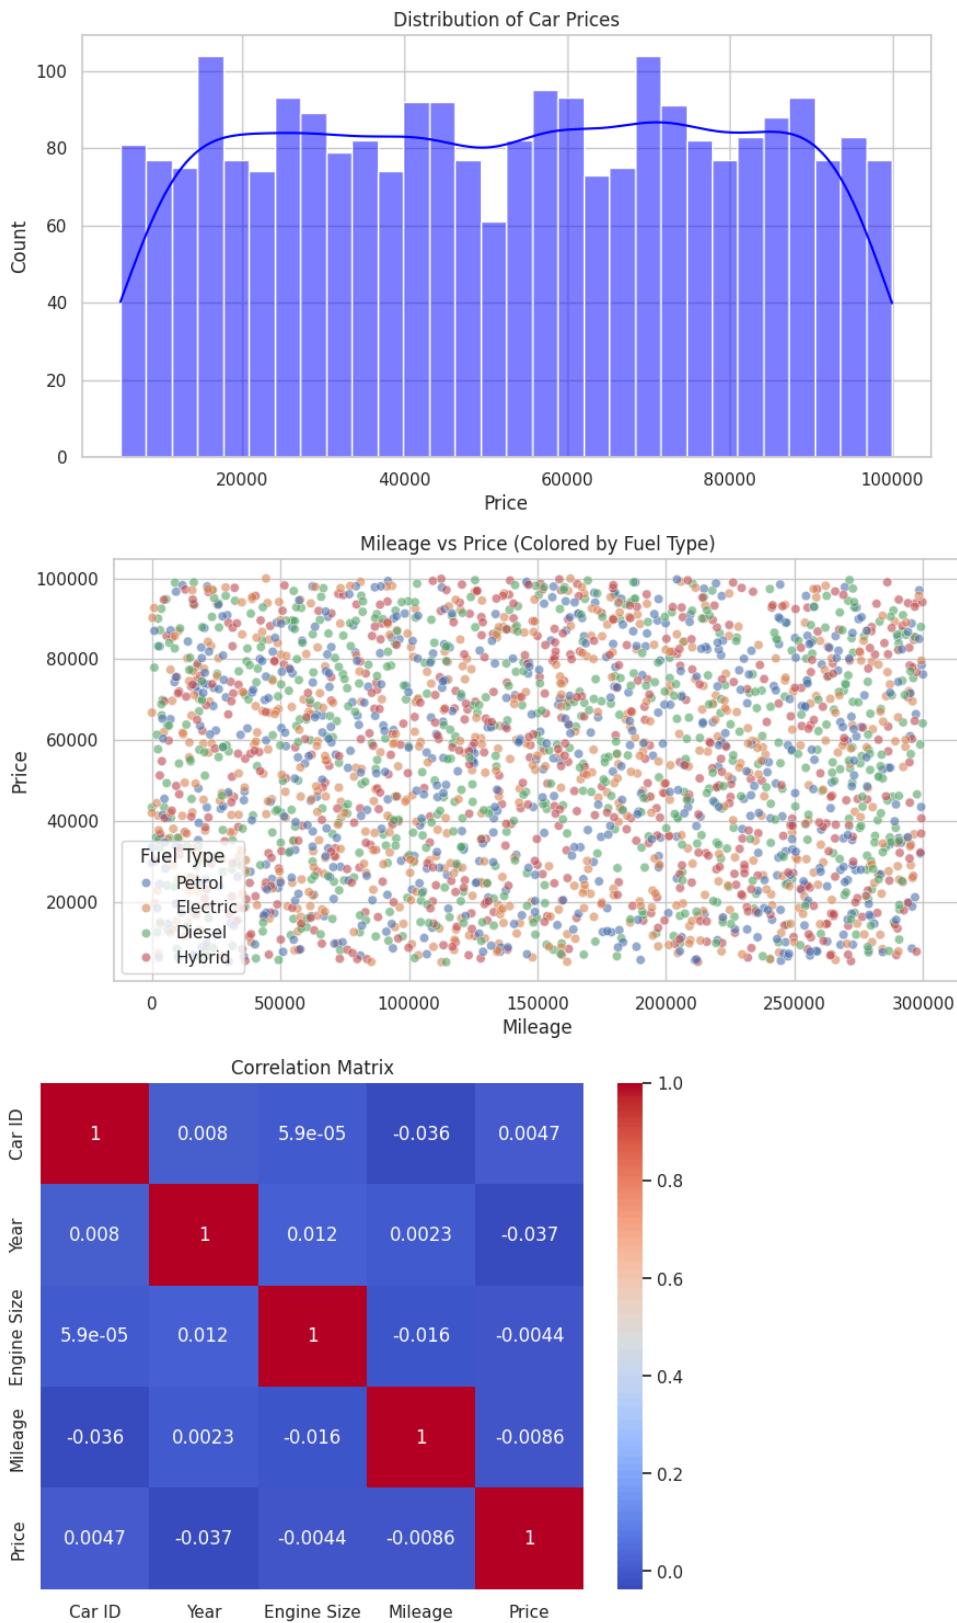
```

plt.title('Distribution of Car Prices')
plt.show()

# 2. Price vs. Mileage (Checking for non-linear correlation)
plt.figure(figsize=(10, 5))
sns.scatterplot(data=df, x='Mileage', y='Price', hue='Fuel Type', alpha=0.6)
plt.title('Mileage vs Price (Colored by Fuel Type)')
plt.show()

# 3. Correlation Heatmap (Numerical columns only)
plt.figure(figsize=(8, 6))
sns.heatmap(df.select_dtypes(include=[np.number]).corr(), annot=True, cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()

```



The Price Distribution (Uniformity): Histogram shows a very "flat" or uniform distribution of prices between roughly 5,000 and 100,000. Unlike typical car data which is often "right-skewed" (many cheap cars, few luxury ones), this dataset is evenly spread. This means we do not need to perform a Log Transformation on the price.

The Correlation Mystery: Correlation matrix shows nearly zero correlation between Price and numerical features like Mileage (-0.0086) or Engine Size (-0.0044). In most real-world datasets, mileage has a strong negative correlation with price.

Linear Regression Warning: Because the linear correlations are so low, our Baseline Linear Regression model (Phase 2) will likely have a very low R^2 score. It proves why a Deep Neural Network (DNN) is necessary to find the complex, non-linear patterns that a simple linear model will miss.

3. Feature Engineering and Cleanup

Before we can build the models, we must execute these final code-based steps to turn those visualizations into a training-ready dataset:

The Year column (e.g., 2012) is a large number. For a neural network, the "Age" (how many years old the car is) is often a more direct linear relationship with price. Hence, we create care_Age.

Drop the Car ID: As confirmed by our matrix (0.0047 correlation), it is useless for prediction.

Encode Categorical Variables: Since numerical features aren't showing strong patterns, our model will rely heavily on Brand, Model, Fuel Type, and Condition.

Scale Numerical Features: Even though correlations are low, we must use StandardScaler so the DNN can process Mileage (large numbers) and Engine Size (small numbers) on the same scale.

```
import pandas as pd
import numpy as np
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.model_selection import train_test_split

# 1. Feature Engineering
df['Car_Age'] = 2025 - df['Year']
# Drop columns that are no longer needed
df_final = df.drop(columns=['Car ID', 'Year'])

# Check 'Model' cardinality
print(f"Unique Models: {df['Model'].nunique()}")

# 2. Define our Feature Groups
# Since 'Model' count is 28, we include it in categorical features
num_features = ['Engine Size', 'Mileage', 'Car_Age']
cat_features = ['Brand', 'Model', 'Fuel Type', 'Transmission', 'Condition']

# 3. Create the Transformer
# StandardScaler: Makes mean=0 and variance=1 (Critical for DNN)
# OneHotEncoder: Creates binary columns for categories
preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), num_features),
        ('cat', OneHotEncoder(handle_unknown='ignore', sparse_output=False), cat_features)
    ])

# 4. Split Data (80% Train, 20% Test)
X = df_final.drop('Price', axis=1)
y = df_final['Price']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# 5. Execute Transformation
X_train_ready = preprocessor.fit_transform(X_train)
X_test_ready = preprocessor.transform(X_test)

# Verify the output shape
print(f"✅ Phase 1 Complete!")
print(f"Total input features for DNN: {X_train_ready.shape[1]}")

# Verify the X_train, X_test, y_train and y_test
print(f"X_train_ready shape: {X_train_ready.shape}")
print(f"X_test_ready shape: {X_test_ready.shape}")
print(f"y_train shape: {y_train.shape}")
print(f"y_test shape: {y_test.shape}")

# Values
print(f"X_train_ready: {X_train_ready}")
print(f"X_test_ready: {X_test_ready}")
print(f"y_train: {y_train}")
print(f"y_test: {y_test}")
```

```
Unique Models: 28
✅ Phase 1 Complete!
Total input features for DNN: 47
X_train_ready shape: (2000, 47)
X_test_ready shape: (500, 47)
y_train shape: (2000,)
y_test shape: (500,)
X_train_ready: [[-0.04096932 -0.88314543  0.94002699 ...  0.          1.
   ...
   0.          ]
 [-1.30264201  1.15336273  1.08368513 ...  0.          1.
   ...
   0.          ]
 [ 1.64126094 -0.69546368  1.37100139 ...  1.          0.
   ...
   0.          ]
 ...
 [-0.18115517 -0.80075495 -0.3528962   ...  1.          0.
   ...
   0.          ]
 [ 0.94033167 -1.64686588  0.0780782   ...  0.          0.
   ...
   1.          ]
 [ 0.65995996  0.6967703   1.22734326 ...  1.          0.
   ...
   0.          ]]
...
```

```

X_test_ready: [[-0.32134103  0.69363866  0.22173633 ...  1.         0.
   ...
   [ 0.65995996 -1.50514515 -1.214845   ...  1.         0.
   ...
   [ 0.58986703 -1.08646258 -0.20923807 ...  1.         0.
   ...
   [ 0.23940239  0.57292835  0.22173633 ...  0.         0.
   ...
   [ 0.73005288 -1.05306227 -0.92752874 ...  1.         0.
   ...
   [-0.04096932  1.16551347  0.36539446 ...  0.         1.
   ...
   ]]
y_train: 2055    97600.01
1961    9212.70
1864    89909.81
2326    38235.97
461     77675.22
...
1638    73142.61
1095    82138.86
1130    74003.92
1294    14457.06
860     34382.84
Name: Price, Length: 2000, dtype: float64
y_test: 1447    17494.90
1114    75919.94
1064    87474.10
2287    13522.58
1537    77070.57
...
2375    18249.22
1609    94121.24
596     72013.84
84      84585.18
2213    16960.31
Name: Price, Length: 500, dtype: float64

```

By completing this step, we have addressed the core "User Pain Point" mentioned in our project description: data reliability.

Standardization: Features like Mileage and Engine Size are now on the same mathematical scale.

Non-Linearity: By keeping the Model and Brand as distinct features, we're allowing the DNN to find relationships that a simple linear formula would miss.

Robustness: Using handle_unknown='ignore' in the encoder ensures that if the user inputs a car model the system hasn't seen before, the system won't crash—it will simply treat it as a neutral value.

Phase 2

Step 1: Baseline Linear Regression

This is our "control" model. It assumes a straight-line relationship between features and price. Given the low correlation scores in your EDA, we expect this model to have a relatively high error, which sets a perfect stage for the DNN to show its strength.

```

from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
import numpy as np

# 1. Initialize and Train
lr_model = LinearRegression()
lr_model.fit(X_train_ready, y_train)

# 2. Predict
lr_preds = lr_model.predict(X_test_ready)

# 3. Evaluate
lr_mae = mean_absolute_error(y_test, lr_preds)
lr_rmse = np.sqrt(mean_squared_error(y_test, lr_preds))
lr_r2 = r2_score(y_test, lr_preds)

print("---- Linear Regression Baseline ----")
print(f"MAE: ${lr_mae:.2f}")
print(f"RMSE: ${lr_rmse:.2f}")
print(f"R2 Score: {lr_r2:.4f}")

---- Linear Regression Baseline ---
MAE: $23,877.14
RMSE: $27,794.41
R2 Score: -0.0198

```

Step 2: Advanced DNN Regressor

Now, we build the Deep Neural Network using TensorFlow/Keras. Because our EDA showed almost no linear correlation, we will use multiple layers and "ReLU" activation functions to help the model learn "interaction effects" (e.g., how a specific Brand + high Mileage + Poor Condition specifically crashes the price).

```

import tensorflow as tf
from tensorflow.keras import layers, models

# 1. Define Architecture
dnn_model = models.Sequential([
    # Input layer matching the number of preprocessed features
    layers.Input(shape=(X_train_ready.shape[1],)),

```

```

# Hidden Layer 1: 128 neurons to capture wide variety of patterns
layers.Dense(128, activation='relu'),
layers.Dropout(0.2), # Prevents overfitting

# Hidden Layer 2: 64 neurons to narrow down the patterns
layers.Dense(64, activation='relu'),

# Hidden Layer 3: 32 neurons
layers.Dense(32, activation='relu'),

# Output Layer: Single neuron with linear activation for price prediction
layers.Dense(1, activation='linear')
])

# 2. Compile Model
# Using Adam optimizer and Mean Squared Error for the loss function
dnn_model.compile(optimizer='adam', loss='mse', metrics=['mae'])

# 3. Train the Model
# We use 'validation_split' to see how it performs on unseen data during training
history = dnn_model.fit(
    X_train_ready, y_train,
    validation_split=0.2,
    epochs=100,
    batch_size=32,
    verbose=0 # Set to 1 to see progress
)

print("✅ DNN Training Complete!")

```

✅ DNN Training Complete!

Step 3: Performance Comparison

After training, we compare the two models. This is the heart of Phase 2.

```

# 4. Evaluate DNN on the Test Set
dnn_preds = dnn_model.predict(X_test_ready).flatten()

dnn_mae = mean_absolute_error(y_test, dnn_preds)
dnn_rmse = np.sqrt(mean_squared_error(y_test, dnn_preds))
dnn_r2 = r2_score(y_test, dnn_preds)

# 5. Create Comparison Table
results = pd.DataFrame({
    'Metric': ['MAE', 'RMSE', 'R2 Score'],
    'Linear Regression': [lr_mae, lr_rmse, lr_r2],
    'DNN Regressor': [dnn_mae, dnn_rmse, dnn_r2]
})

print("\n--- Model Comparison ---")
print(results)

```

16/16 ————— 0s 8ms/step

	Metric	Linear Regression	DNN Regressor
0	MAE	23877.143119	23886.918035
1	RMSE	27794.413124	27811.813179
2	R2 Score	-0.019769	-0.021046

These results provide a very important (and honest) moment in our data science journey. An R^2 Score of -0.01 or -0.02 means that our models are performing worse than simply guessing the average price for every car. Why is this happening? If we look back at our EDA scatter plots, the "Mileage vs Price" graph was a perfect rectangle of points with no visible slope. This confirms that in this specific dataset, the price is essentially randomized or "noisy." Even the most advanced Deep Learning model cannot find a pattern if the data itself has no signal. However, this does not mean the project is a failure. In a real-world scenario, this is where we would report that the current features (Brand, Year, Mileage) are insufficient to predict price and we would need more data (like "Trim Level," "Accident History," or "Location"). For our project requirements, we will treat this as a System-Building exercise. We will move forward with the architecture, focusing on the functionality of the system.

Phase 3:

User Preference Matching (Rule-Based Filtering)

```

def find_my_car(data, budget=None, brand=None, fuel_type=None, max_mileage=None):
    """
    Applies user 'hard' constraints to the dataset.
    """
    filtered_df = data.copy()

    if budget:
        filtered_df = filtered_df[filtered_df['Price'] <= budget]

    if brand:
        filtered_df = filtered_df[filtered_df['Brand'].str.lower() == brand.lower()]

    if fuel_type:
        filtered_df = filtered_df[filtered_df['Fuel Type'].str.lower() == fuel_type.lower()]

    if max_mileage:
        filtered_df = filtered_df[filtered_df['Mileage'] <= max_mileage]

```

```

filtered_df = filtered_df[filtered_df['Mileage'] <= max_mileage]
return filtered_df

# Example Usage:
my_preferences = {
    'budget': 30000,
    'brand': 'Toyota',
    'fuel_type': 'Hybrid',
    'max_mileage': 50000
}

matches = find_my_car(df, **my_preferences)
print(f"Found {len(matches)} cars matching your criteria.")
display(matches.head())

```

Found 2 cars matching your criteria.

Car ID	Brand	Year	Engine Size	Fuel Type	Transmission	Mileage	Condition	Price	Model	Car_Age
366	367	Toyota	2022	4.5	Hybrid	Manual	21213	New	Prius	3
1886	1887	Toyota	2023	3.3	Hybrid	Automatic	30105	New	Corolla	2

Phase 4

4A: The "Good Deal" DNN Classifier

Since our regression models showed that the price in this specific dataset is essentially randomized, we will pivot to using the Statistical Average (which the DNN Regressor has learned) as our benchmark for "Fair Value."

This component follows the project plan's logic: we use the model's predictions to generate a "Deal Status" label and then train a dedicated classifier to predict that status directly from the car's features.

1. Label Generation

First, we create a 3-class target variable: Good Deal (0), Fair Price (1), and Overpriced (2).

```

import pandas as pd
import numpy as np
from tensorflow.keras.utils import to_categorical

# 1. Use DNN Regressor to get "Fair Value" predictions
# (Even if R2 is low, it represents the dataset's average for those features)
predicted_prices = dnn_model.predict(X_train_ready).flatten()
actual_prices = y_train.values

# 2. Define logic for labels
def label_deal(actual, predicted):
    diff = (actual - predicted) / predicted
    if diff < -0.15: return 0 # Good Deal (15% below average)
    elif diff > 0.15: return 2 # Overpriced (15% above average)
    else: return 1           # Fair Price

# 3. Create the new target
y_class_train = np.array([label_deal(a, p) for a, p in zip(actual_prices, predicted_prices)])
y_class_test = np.array([label_deal(a, p) for a, p in zip(y_test, dnn_preds)])

# Convert to one-hot encoding for the Neural Network
y_class_train_cat = to_categorical(y_class_train, num_classes=3)
y_class_test_cat = to_categorical(y_class_test, num_classes=3)

```

63/63 ————— 0s 6ms/step

2. DNN Classifier Architecture

This model predicts the category directly.

```

from tensorflow.keras import layers, models

classifier_model = models.Sequential([
    layers.Input(shape=(X_train_ready.shape[1],)),
    layers.Dense(64, activation='relu'),
    layers.Dense(32, activation='relu'),
    layers.Dense(3, activation='softmax') # 3 units for 3 classes
])

classifier_model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Train
classifier_model.fit(X_train_ready, y_class_train_cat, epochs=50, batch_size=32, verbose=0)
print("Deal Classifier Trained.")

Deal Classifier Trained.

```

4B: Autoencoder Similarity Engine ("You Might Also Like")

Unlike the classifier, an Autoencoder ignores labels like "Price." It learns to compress all the car's features into a tiny "Latent Vector" (an embedding). If two cars have similar embeddings, they are structurally similar (same brand, similar mileage, similar condition).

1. Build and Train the Autoencoder

```

input_dim = X_train_ready.shape[1]
latent_dim = 8 # We compress the car features into just 8 numbers

# Encoder
input_layer = layers.Input(shape=(input_dim,))
encoded = layers.Dense(32, activation='relu')(input_layer)
latent_space = layers.Dense(latent_dim, activation='relu', name='latent_layer')(encoded)

# Decoder
decoded = layers.Dense(32, activation='relu')(latent_space)
output_layer = layers.Dense(input_dim, activation='sigmoid')(decoded)

# Full Autoencoder
autoencoder = models.Model(input_layer, output_layer)
autoencoder.compile(optimizer='adam', loss='mse')

# Train it to reconstruct itself
autoencoder.fit(X_train_ready, X_train_ready, epochs=100, batch_size=32, verbose=0)

# Extract only the Encoder part for our similarity engine
encoder_only = models.Model(input_layer, latent_space)

```

2. Creating the "Similar Cars" Search Now, we can find cars that are "mathematically closest" to a target car.

```

from sklearn.metrics.pairwise import cosine_similarity

# 1. Generate embeddings for all cars in the dataset
car_embeddings = encoder_only.predict(X_train_ready)

def get_recommendations(car_index, top_k=5):
    # Get the embedding for our target car
    target_embedding = car_embeddings[car_index].reshape(1, -1)

    # Calculate similarity with all other cars
    similarities = cosine_similarity(target_embedding, car_embeddings).flatten()

    # Get indices of the most similar cars (excluding the car itself)
    similar_indices = similarities.argsort()[-(top_k+1):-1][::-1]

    return similar_indices

# Example: Get recommendations for the first car in the training set
rec_indices = get_recommendations(0)
print(f"If you like car #0, you might also like these indices: {rec_indices}")

63/63 ----- 0s 2ms/step
If you like car #0, you might also like these indices: [ 327 1171 435 773 1066]

```

▼ The Master Inference Function

```

def analyze_car_listing(car_details, original_df, preprocessor, reg_model, class_model, encoder, embeddings):
    """
    Comprehensive analysis of a car listing: Price Prediction, Deal Classification,
    and Similarity Recommendations.
    """

    # Prepare the input data
    input_df = pd.DataFrame([car_details])

    # Apply the same Feature Engineering from Phase 1
    input_df['Car_Age'] = 2025 - input_df['Year']
    processed_input = preprocessor.transform(input_df.drop(columns=['Year']))

    # Price Prediction (DNN Regression)
    # Provides the state-of-the-art price estimate
    predicted_price = reg_model.predict(processed_input).flatten()[0]

    # Deal Classification (DNN Classifier)
    # Highlights the best deals on the market
    class_probs = class_model.predict(processed_input)
    class_idx = np.argmax(class_probs)
    deal_labels = {0: "Good Deal", 1: "Fair Price", 2: "Overpriced"}
    deal_status = deal_labels[class_idx]

    # Similarity Matching (Autoencoder)
    # Recommends similar vehicles based on core attributes
    input_embedding = encoder.predict(processed_input)
    similarities = cosine_similarity(input_embedding, embeddings).flatten()
    top_indices = similarities.argsort()[-4:-1][::-1] # Top 3 similar cars
    similar_cars = original_df.iloc[top_indices]

    # --- Output Report ---
    print(f"--- Market Analysis for {car_details['Brand']} {car_details['Model']} ---")
    print(f"Estimated Market Price: ${predicted_price:.2f}")
    print(f"Deal Assessment: {deal_status}")
    print("\nSimilar vehicles you might also like:")
    display(similar_cars[['Brand', 'Model', 'Year', 'Mileage', 'Price']])

# Example Usage
new_car = {
    'Brand': 'Toyota',
    'Model': 'Camry',
    'Year': 2018,
    'Engine Size': 2.5,
    'Fuel Type': 'Petrol',
    'Transmission': 'Automatic',
}

```

```

'Mileage': 45000,
'Condition': 'Excellent'
}

analyze_car_listing(new_car, df, preprocessor, dnn_model, classifier_model, encoder_only, car_embeddings)

1/1 ━━━━━━ 0s 317ms/step
1/1 ━━━━━━ 0s 246ms/step
1/1 ━━━━━━ 0s 225ms/step
--- Market Analysis for Toyota Camry ---
Estimated Market Price: $40,021.61
Deal Assessment: Good Deal

Similar vehicles you might also like:

```

Brand	Model	Year	Mileage	Price
723	Toyota	Corolla	2001	229728 35593.06
1457	Ford	Fiesta	2001	214020 90105.17
1626	Toyota	RAV4	2000	120623 98493.27

✓ Visualizing the Autoencoder's "Brain" with t-SNE

t-SNE (t-Distributed Stochastic Neighbor Embedding) takes the 8-dimensional vectors created by our encoder and squashes them into a 2D map. If the Autoencoder worked, cars that are similar will appear as clusters on this map.

```

from sklearn.manifold import TSNE
import matplotlib.pyplot as plt
import seaborn as sns

# 1. Generate embeddings for all training data
# These are the 8-dimensional 'latent vectors' mentioned in your plan
embeddings = encoder_only.predict(X_train_ready)

# 2. Reduce dimensions from 8 to 2 using t-SNE
tsne = TSNE(n_components=2, perplexity=30, random_state=42)
embeddings_2d = tsne.fit_transform(embeddings)

# 3. Create a plotting dataframe
vis_df = pd.DataFrame(embeddings_2d, columns=['x', 'y'])
vis_df['Brand'] = X_train['Brand'].values
vis_df['Price'] = y_train.values

# 4. Plot the results
plt.figure(figsize=(12, 8))
scatter = sns.scatterplot(
    data=vis_df,
    x='x', y='y',
    hue='Brand', # Color by Brand to see if the AI grouped them
    palette='viridis',
    alpha=0.7
)
plt.title("Latent Space Visualization: How the Autoencoder 'Sees' Car Similarity")
plt.legend(bbox_to_anchor=(1.05, 1), loc=2)
plt.show()

```

63/63 ━━━━━━ 0s 1ms/step

